

An R Markdown document converted from Lab 8

Lab 8:

- Decision Tree
- Random Forest

```
library(tree)
```

```
library(repr)
options(repr.plot.width=12, repr.plot.height=6)
```

Classification Tree

- Use the `Carseats` dataset
- Transform a continuous variable `Sales` to a binary variable.

```
library(ISLR2)
attach(Carseats)
```

```
High <- factor(ifelse(Sales <= 8, "No", "Yes"))
```

```
Carseats <- data.frame(Carseats, High)
```

Use `tree` function to fit a classification tree to predict `High` using all variables but `Sales`.

```
tree.carseats <- tree(High ~ . - Sales, Carseats)
```

```
head(Carseats)
```

```
##   Sales CompPrice Income Advertising Population Price ShelveLoc Age Education
## 1  9.50      138     73          11        276   120         Bad  42         17
## 2 11.22      111     48          16        260    83         Good  65         10
## 3 10.06      113     35          10        269    80       Medium  59         12
## 4  7.40      117    100           4        466    97       Medium  55         14
## 5  4.15      141     64           3        340   128         Bad  38         13
## 6 10.81      124    113          13        501    72         Bad  78         16
##   Urban  US High
## 1  Yes Yes  Yes
## 2  Yes Yes  Yes
## 3  Yes Yes  Yes
## 4  Yes Yes   No
## 5  Yes No   No
## 6   No Yes  Yes
```

```
summary(tree.carseats)
```

```
##
## Classification tree:
## tree(formula = High ~ . - Sales, data = Carseats)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Income" "CompPrice" "Population"
## [6] "Advertising" "Age" "US"
## Number of terminal nodes: 27
## Residual mean deviance: 0.4575 = 170.7 / 373
## Misclassification error rate: 0.09 = 36 / 400
```

The deviance is given by

$$-2 \sum_m \sum_k n_{mk} \log \hat{p}_{mk}$$

where n_{mk} is the number of observations in the m th terminal node that belong to the k th class and \hat{p}_{mk} represents the proportion of training observations in the m -th region that are from the k th class.

It is closely related to the entropy

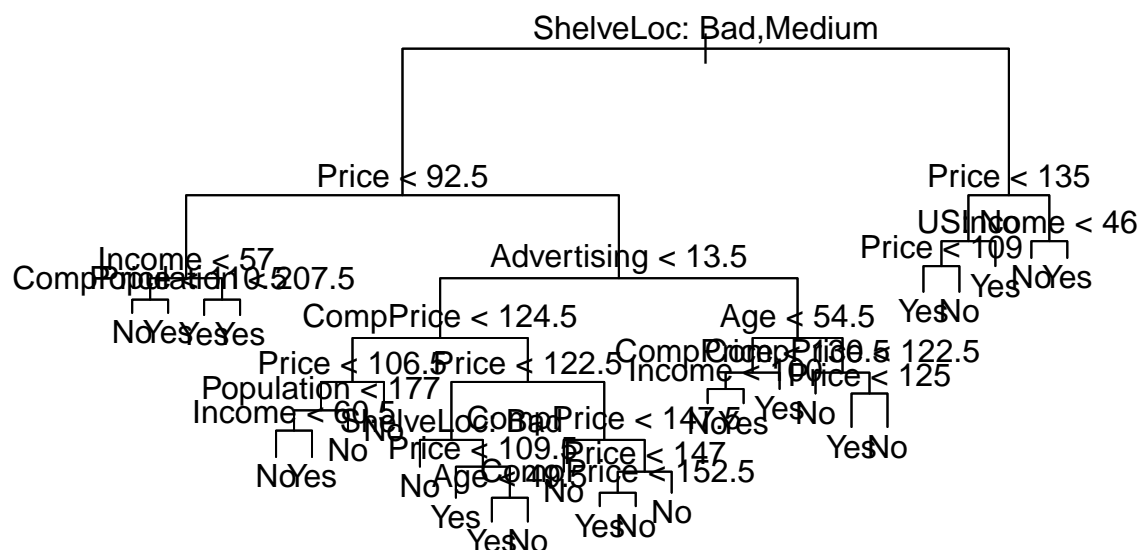
$$-\sum_m \sum_k \hat{p}_{mk} \log \hat{p}_{mk}$$

A small deviance indicates a tree provide a good fit to the training data.

The *residual mean deviance* is simply the deviance divided by $n - |T_0|$ and in this case $400 - 27 = 373$.

We can also plot a tree with plot function.

```
plot(tree.carseats)
text(tree.carseats, pretty = 0) # pretty means that to include the category names for any qualitative v
```



Seems like that the most important variable is shelving location, as the first branch differentiates Good locations from Bad and Medium.

It turns out that the left branch corresponds to the split being true. So the left branch is for Bad and Medium, while the right branch is for Good.

```
tree.carseats
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
##      1) root 400 541.500 No ( 0.59000 0.41000 )
##      2) ShelfLoc: Bad,Medium 315 390.600 No ( 0.68889 0.31111 )
##      4) Price < 92.5 46 56.530 Yes ( 0.30435 0.69565 )
##      8) Income < 57 10 12.220 No ( 0.70000 0.30000 )
##      16) CompPrice < 110.5 5 0.000 No ( 1.00000 0.00000 ) *
##      17) CompPrice > 110.5 5 6.730 Yes ( 0.40000 0.60000 ) *
##      9) Income > 57 36 35.470 Yes ( 0.19444 0.80556 )
##      18) Population < 207.5 16 21.170 Yes ( 0.37500 0.62500 ) *
##      19) Population > 207.5 20 7.941 Yes ( 0.05000 0.95000 ) *
##      5) Price > 92.5 269 299.800 No ( 0.75465 0.24535 )
##      10) Advertising < 13.5 224 213.200 No ( 0.81696 0.18304 )
##      20) CompPrice < 124.5 96 44.890 No ( 0.93750 0.06250 )
##      40) Price < 106.5 38 33.150 No ( 0.84211 0.15789 )
##      80) Population < 177 12 16.300 No ( 0.58333 0.41667 )
##      160) Income < 60.5 6 0.000 No ( 1.00000 0.00000 ) *
##      161) Income > 60.5 6 5.407 Yes ( 0.16667 0.83333 ) *
##      81) Population > 177 26 8.477 No ( 0.96154 0.03846 ) *
##      41) Price > 106.5 58 0.000 No ( 1.00000 0.00000 ) *
##      21) CompPrice > 124.5 128 150.200 No ( 0.72656 0.27344 )
##      42) Price < 122.5 51 70.680 Yes ( 0.49020 0.50980 )
##      84) ShelfLoc: Bad 11 6.702 No ( 0.90909 0.09091 ) *
##      85) ShelfLoc: Medium 40 52.930 Yes ( 0.37500 0.62500 )
##      170) Price < 109.5 16 7.481 Yes ( 0.06250 0.93750 ) *
##      171) Price > 109.5 24 32.600 No ( 0.58333 0.41667 )
##      342) Age < 49.5 13 16.050 Yes ( 0.30769 0.69231 ) *
##      343) Age > 49.5 11 6.702 No ( 0.90909 0.09091 ) *
##      43) Price > 122.5 77 55.540 No ( 0.88312 0.11688 )
##      86) CompPrice < 147.5 58 17.400 No ( 0.96552 0.03448 ) *
##      87) CompPrice > 147.5 19 25.010 No ( 0.63158 0.36842 )
##      174) Price < 147 12 16.300 Yes ( 0.41667 0.58333 )
##      348) CompPrice < 152.5 7 5.742 Yes ( 0.14286 0.85714 ) *
##      349) CompPrice > 152.5 5 5.004 No ( 0.80000 0.20000 ) *
##      175) Price > 147 7 0.000 No ( 1.00000 0.00000 ) *
##      11) Advertising > 13.5 45 61.830 Yes ( 0.44444 0.55556 )
##      22) Age < 54.5 25 25.020 Yes ( 0.20000 0.80000 )
##      44) CompPrice < 130.5 14 18.250 Yes ( 0.35714 0.64286 )
##      88) Income < 100 9 12.370 No ( 0.55556 0.44444 ) *
##      89) Income > 100 5 0.000 Yes ( 0.00000 1.00000 ) *
##      45) CompPrice > 130.5 11 0.000 Yes ( 0.00000 1.00000 ) *
##      23) Age > 54.5 20 22.490 No ( 0.75000 0.25000 )
##      46) CompPrice < 122.5 10 0.000 No ( 1.00000 0.00000 ) *
##      47) CompPrice > 122.5 10 13.860 No ( 0.50000 0.50000 )
```

```
##           94) Price < 125 5    0.000 Yes ( 0.00000 1.00000 ) *
##           95) Price > 125 5    0.000 No  ( 1.00000 0.00000 ) *
##      3) ShelfLoc: Good 85   90.330 Yes ( 0.22353 0.77647 )
##           6) Price < 135 68   49.260 Yes ( 0.11765 0.88235 )
##           12) US: No 17    22.070 Yes ( 0.35294 0.64706 )
##           24) Price < 109 8    0.000 Yes ( 0.00000 1.00000 ) *
##           25) Price > 109 9   11.460 No  ( 0.66667 0.33333 ) *
##           13) US: Yes 51   16.880 Yes ( 0.03922 0.96078 ) *
##           7) Price > 135 17   22.070 No  ( 0.64706 0.35294 )
##           14) Income < 46 6    0.000 No  ( 1.00000 0.00000 ) *
##           15) Income > 46 11   15.160 Yes ( 0.45455 0.54545 ) *
```

Now estimate test error

```
set.seed(2)
train <- sample(1 : nrow(Carseats), 200)
Carseats.test <- Carseats[-train, ]
High.test <- High[-train]

tree.carseats <- tree(High ~ . - Sales, Carseats, subset=train)

# We specify type = class to instruct R to return the actual class prediction.
tree.pred <- predict(tree.carseats, Carseats.test, type="class")
table(tree.pred, High.test)
```

```
##           High.test
## tree.pred  No Yes
##           No 104 33
##           Yes 13 50
```

```
(104 + 50) / 200
```

```
## [1] 0.77
```

Use `cv.tree` to perform cross validation

Cost complexity pruning is used in order to select a sequence of trees for consideration

- Estimating the cross-validation error for every possible subtree would be computationally impossible, given that there are so many possible subtrees
- We start with a very big tree T_0 , then we consider a sequence of trees indexed by a nonnegative tuning parameter α .
- For each value of α , there corresponds to a subtree $T \subset T_0$ such that

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

is as small as possible.

- $|T|$ is the number of terminal nodes, R_m is the rectangle for the m -th terminal node, \hat{y}_{R_m} is the predicted response for R_m .

- As we increase α from 0, branches get pruned from the tree in a nested and predictable fashion, so obtaining the whole sequence of subtrees as a function of α is easy.
- Then we can use cross validation to select among this sequence of trees. ### We use argument `FUN = prune.misclass` in order to indicate that we want to use classification error as the metric to guide the cross-validation and pruning process ### The default for `cv.tree` is deviance.

```
set.seed(7)
cv.carseats <- cv.tree(tree.carseats, FUN=prune.misclass)
names(cv.carseats)
```

```
## [1] "size" "dev" "k" "method"
```

Size stands for the number of terminal nodes and k is the cost-complexity parameter (corresponds to α)

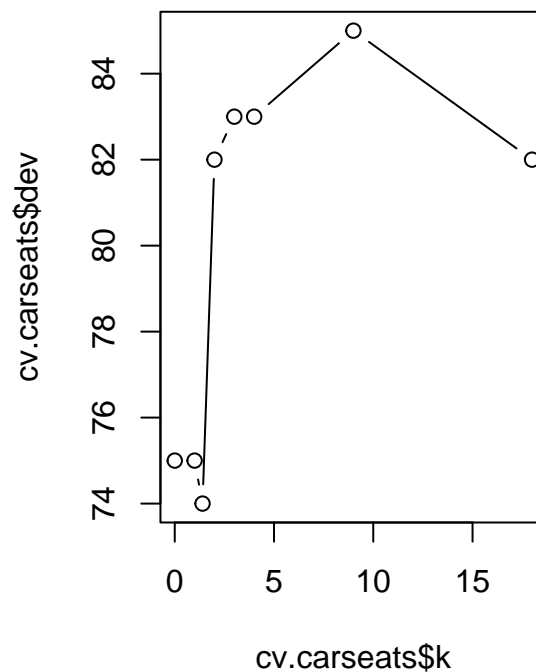
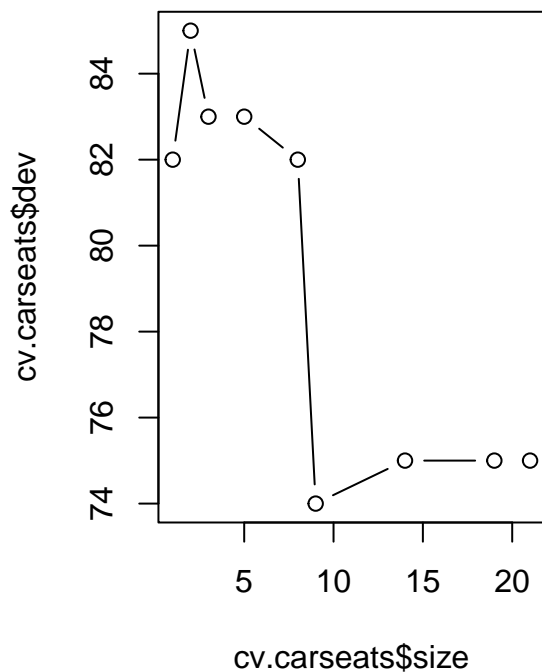
```
cv.carseats

## $size
## [1] 21 19 14 9 8 5 3 2 1
##
## $dev
## [1] 75 75 75 74 82 83 83 85 82
##
## $k
## [1] -Inf 0.0 1.0 1.4 2.0 3.0 4.0 9.0 18.0
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

Dev actually means the misclassification errors

We should then select the tree with 9 terminal nodes

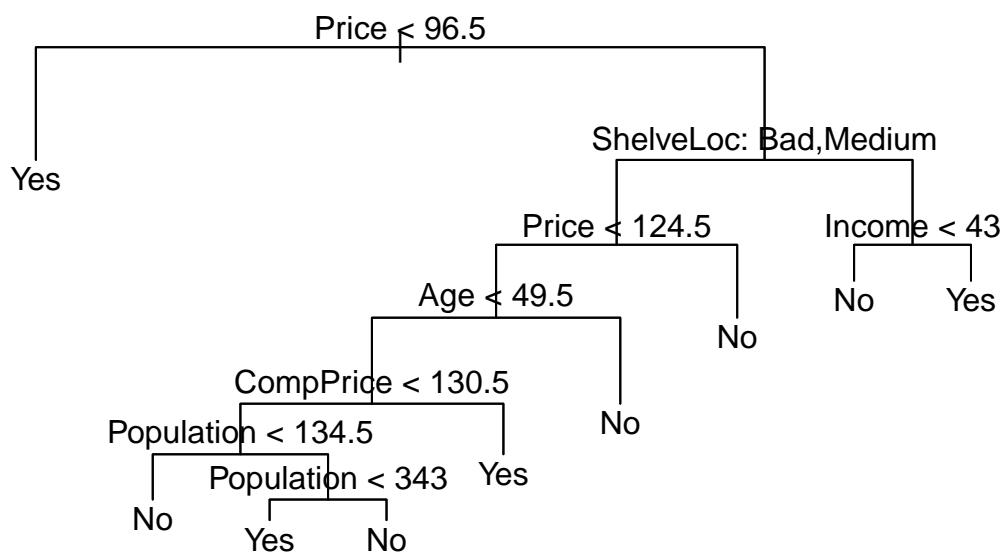
```
par(mfrow = c(1, 2))
plot(cv.carseats$size, cv.carseats$dev, type="b")
plot(cv.carseats$k, cv.carseats$dev, type="b")
```



Use `prune.misclass` function to prune the tree to obtain the nine-node tree

```
prune.carseats <- prune.misclass(tree.carseats, best=9)
```

```
plot(prune.carseats)
text(prune.carseats, pretty=0)
```



```
tree.pred <- predict(prune.carseats, Carseats.test, type="class")
table(tree.pred, High.test)
```

```
##           High.test
```

```
## tree.pred No Yes
##      No  97  25
##      Yes 20  58
```

```
(97 + 58) / 200
```

```
## [1] 0.775
```

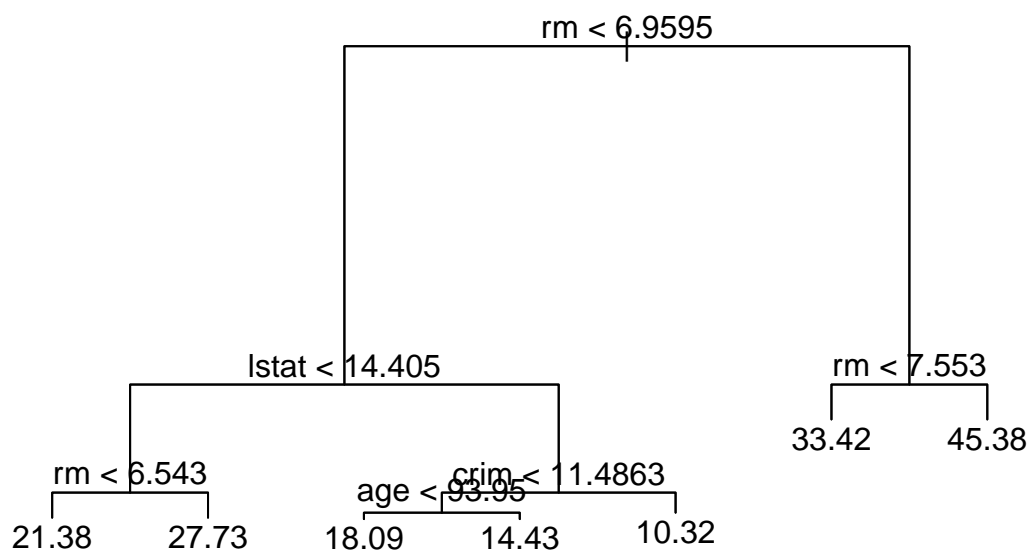
Regression Tree

Now use Boston dataset

```
set.seed(1)
train <- sample(1 : nrow(Boston), nrow(Boston) / 2)
tree.boston <- tree(medv ~ ., Boston, subset = train)
summary(tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm"      "lstat"   "crim"    "age"
## Number of terminal nodes: 7
## Residual mean deviance: 10.38 = 2555 / 246
## Distribution of residuals:
##      Min. 1st Qu.  Median     Mean 3rd Qu.     Max.
## -10.1800 -1.7770 -0.1775  0.0000  1.9230 16.5800
```

```
plot(tree.boston)
text(tree.boston, pretty = 0)
```

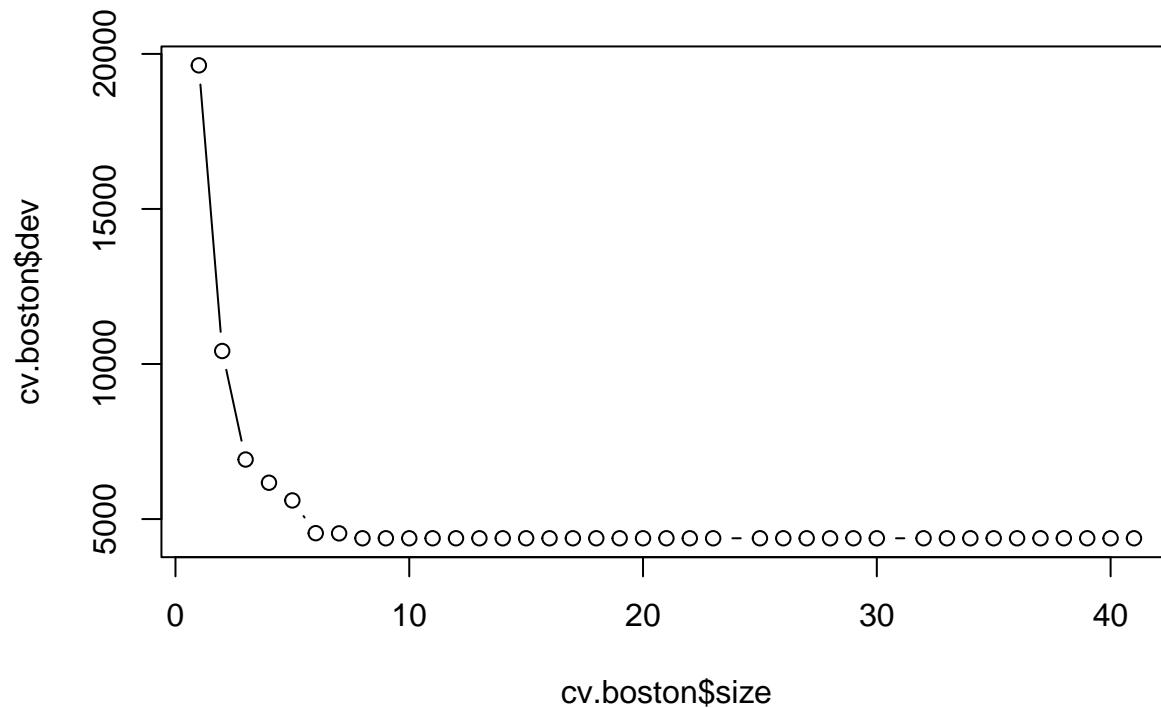


```
tree.boston <- tree(medv ~ ., Boston, subset = train, control = tree.control(nobs = length(train), mind
```

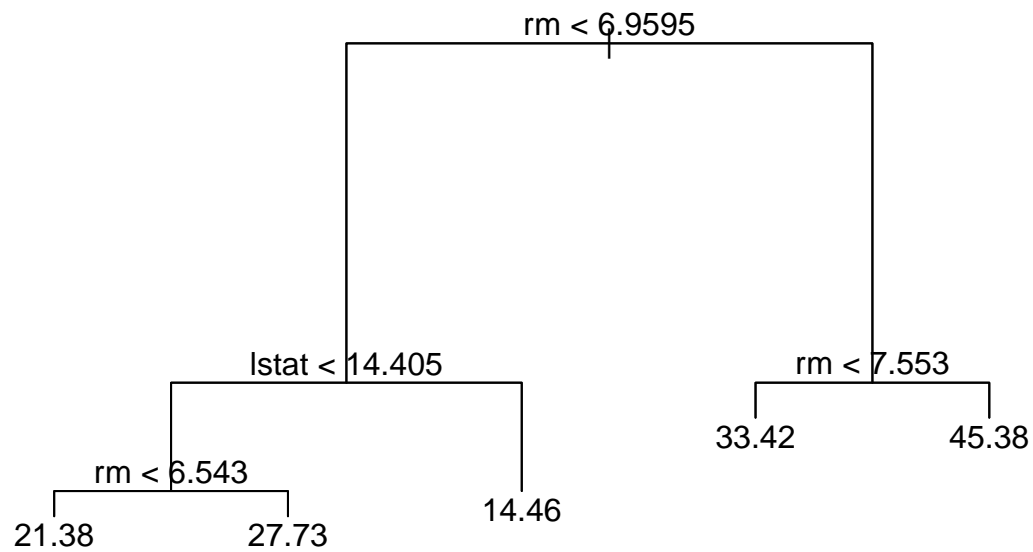
```
summary(tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train, control = tree.control(nobs = length(train),
##   mindev = 0))
## Variables actually used in tree construction:
## [1] "rm"      "lstat"   "indus"   "age"     "nox"     "dis"     "ptratio"
## [8] "tax"     "crim"
## Number of terminal nodes: 41
## Residual mean deviance: 5.542 = 1175 / 212
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -8.140 -1.200   0.000   0.000  1.087  12.860
```

```
cv.boston <- cv.tree(tree.boston)
plot(cv.boston$size, cv.boston$dev, type="b")
```

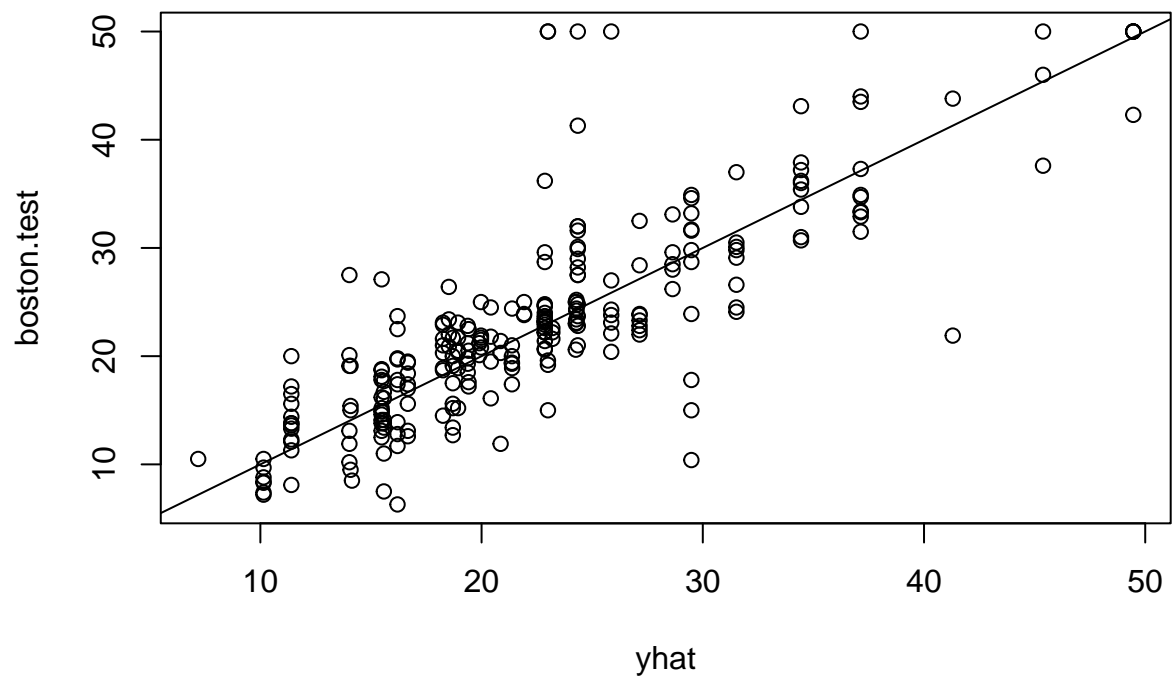


```
prune.boston <- prune.tree(tree.boston, best = 5)
plot(prune.boston)
text(prune.boston, pretty = 0)
```

```
yhat <- predict(tree.boston, newdata = Boston[-train, ])
boston.test <- Boston[-train, "medv"]
```

```
plot(yhat, boston.test)
abline(0, 1)
```



```
mean((yhat - boston.test)^2)
```

```
## [1] 30.19368
```

Bagging and Random Forests

```
library(randomForest)
```

```
## randomForest 4.7-1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

Recal that bagging is simply a special case of a random forest with $m = p$. Therefore, `randomForest` can be used to perform both random forests and bagging.

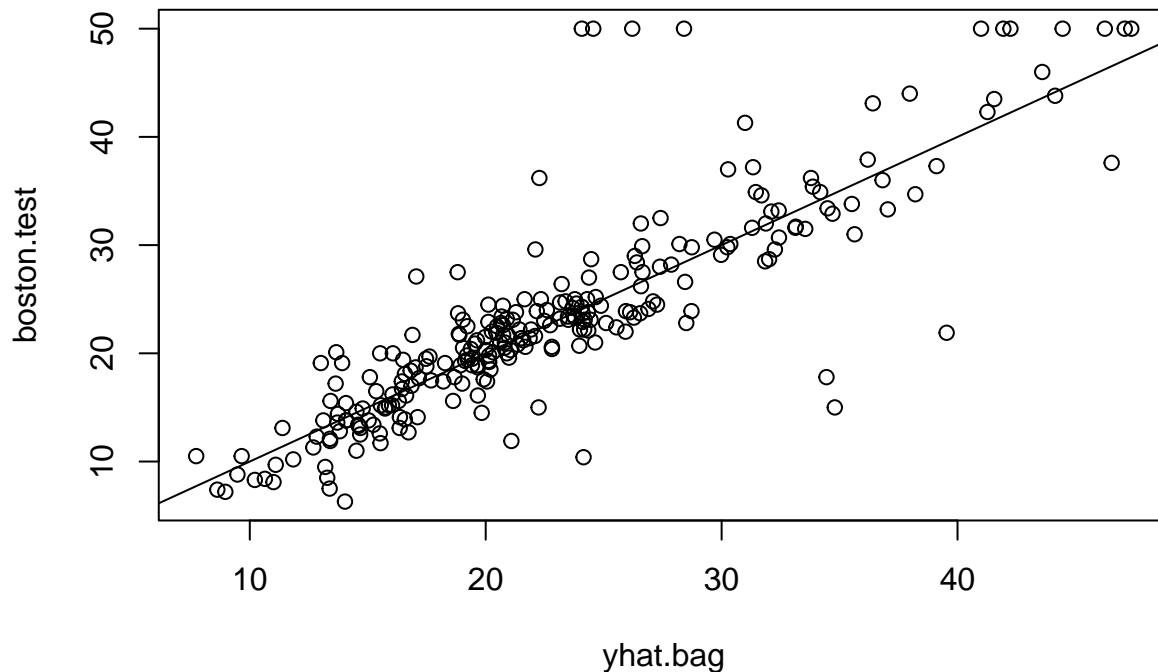
```
set.seed(1)
bag.boston <- randomForest(medv ~ ., data = Boston,
                           subset=train, mtry=12, importance = TRUE)
```

```
bag.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = Boston, mtry = 12, importance = TRUE,      subset = train)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 12
##
##           Mean of squared residuals: 11.40162
##           % Var explained: 85.17
```

```
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])
```

```
plot(yhat.bag, boston.test)
abline(0, 1)
```



```
mean((yhat.bag - boston.test)^2)
```

```
## [1] 23.41916
```

```
bag.boston <- randomForest(medv ~ ., data = Boston,
                           subset = train, mtry=12, ntree=25)
yhat.bag <- predict(bag.boston, newdata = Boston[-train, ])
mean((yhat.bag - boston.test)^2)
```

```
## [1] 25.75055
```

For random forest, we will specify a smaller `mtry` argument. The default is for regression trees, we set $m = p/3$ and for classification trees, we set $m = \sqrt{p}$.

```
set.seed(1)
rf.boston <- randomForest(medv ~ ., data = Boston,
                          subset = train, mtry = 6, importance = T)
yhat.rf <- predict(rf.boston, newdata = Boston[-train, ])
```

```
mean((yhat.rf - boston.test)^2)
```

```
## [1] 20.06644
```

```
importance(rf.boston)
```

```
##           %IncMSE IncNodePurity
## crim      19.435587    1070.42307
```

## zn	3.091630	82.19257
## indus	6.140529	590.09536
## chas	1.370310	36.70356
## nox	13.263466	859.97091
## rm	35.094741	8270.33906
## age	15.144821	634.31220
## dis	9.163776	684.87953
## rad	4.793720	83.18719
## tax	4.410714	292.20949
## ptratio	8.612780	902.20190
## lstat	28.725343	5813.04833

First column is based on the mean decrease of accuracy in predictions on the out-of-bag samples when a given variable is permuted.

The second column is a measure of the total decrease in node impurity that results from splits over that variable, averaged over all trees.

For regression trees, node impurity is measured by the training RSS, and for classification, node impurity is measured by the deviance.

```
varImpPlot(rf.boston)
```

rf.boston

